

Cross-Correlation: AI Availability vs Exploitation Speed vs Breach Disclosures

Analysis Date: 6 June 2026

Purpose: Test whether the "AI is accelerating weaponisation" narrative holds up when you overlay three independent data streams: (1) what AI tools were actually available to threat actors, (2) how fast exploitation was happening, and (3) whether mandated breach disclosures show a corresponding increase.

1. AI Model Availability Timeline — What Actors Could Actually Use

This is the critical control the hype cycle ignores. You can't attribute acceleration to AI if the tools didn't exist yet.

Commercial / Closed Models (require API access, have guardrails)

Date	Model	Offensive Relevance
Nov 2022	ChatGPT (GPT-3.5)	Conversational only. 0% exploit success (Fang et al.)
Mar 2023	GPT-4	First model to autonomously exploit CVEs (87% success, April 2024 research). But: API-gated, logged, guardrailed
Dec 2023	Gemini 1.0	Multimodal but limited coding capability
Mar 2024	Claude 3 (Opus/Sonnet)	Strong coding but heavily safety-aligned
May 2024	GPT-4o	Faster, cheaper, multimodal — lowers cost barrier
Jun 2024	Claude 3.5 Sonnet	SOTA coding benchmark — the model that changed agent capability
Sep 2024	o1 (reasoning)	Chain-of-thought reasoning — step change for complex exploit logic

Date	Model	Offensive Relevance
Jan 2025	DeepSeek R1	Open-weight reasoning model (MIT license) — no guardrails if self-hosted
Feb 2025	Claude 3.5 Haiku / GPT-4.5	Cheaper, faster — commoditises access
Apr 2025	Claude 4 Opus / Sonnet	Major coding + agent capability jump
Aug 2025	GPT-5	Unified reasoning engine
Feb 2026	GPT 5.3-Codex	First to hit OpenAI's "High" cybersecurity threshold
Apr 2026	Claude Mythos Preview	Found zero-days in every major OS. Too dangerous for public release

Open-Weight Models (can be self-hosted, fine-tuned, guardrails removed)

Date	Model	Why It Matters for Threat Actors
Jul 2023	Llama 2 (7B–70B)	First frontier-adjacent open model. 0% exploit success (Fang et al.)
Sep 2023	Mistral 7B	Small, fast, easy to fine-tune. Apache 2.0
Dec 2023	Mixtral 8x7B	MoE architecture — better reasoning, still small enough to self-host
Apr 2024	Llama 3 (8B–70B)	Significant capability jump. Open weight
Jul 2024	Llama 3.1 (405B)	First truly frontier-scale open model
Sep 2024	Qwen 2.5 (72B)	Strong coding, Apache 2.0, popular in China
Jan 2025	DeepSeek V3 + R1	671B MoE, MIT license, reasoning capability. Inflection point for open-weight offensive use
Apr 2025	Llama 4 Scout/ Maverick	Next-gen open models
2025–2026	Qwen 3, GLM-5, Kimi K2	Multiple frontier-class open models under permissive licenses

Key Observation: The Tooling Gap

Before mid-2024, the only model with demonstrated exploit capability was GPT-4 — API-gated, logged, and guardrailed. Threat actors couldn't use it without leaving traces or being rate-limited. Open models (Llama 2, Mistral 7B) scored 0% on exploit benchmarks.

Mid-2024 to early 2025 was the transition period. Claude 3.5 Sonnet and GPT-4o made strong coding available cheaply. Llama 3.1 405B brought near-frontier capability to self-hosted setups. But reasoning and multi-step planning were still weak.

January 2025 (DeepSeek R1) is the real inflection point for threat actors. An MIT-licensed, reasoning-capable, 671B model that anyone can download, fine-tune, and run without guardrails. No API logs. No content filtering. No terms of service. This is when the open-weight ecosystem crossed the threshold from "can write code" to "can reason about exploitation."

The China/Claude Code incident (November 2025) shows the alternative path: jailbreaking commercial models rather than using open-weight ones. Both routes converge on the same capability.

2. Exploitation Speed — The Independent Variable

Now overlay what was actually happening in the field:

Time-to-Exploit Trend (Mandiant M-Trends)

Period	Mean TTE	AI Tools Available
2018–2019	63 days	None
2020–2021	~44 days (est.)	GPT-3 (no exploit capability)
2022	~32 days (est.)	ChatGPT launched Nov 2022 (no exploit capability)
2023	~15 days (est.)	GPT-4 available (87% CVE success proven in lab, Apr 2024)
2024	5 days	GPT-4o, Claude 3.5, Llama 3.1, o1 all available
2025	–7 days	

Period	Mean TTE	AI Tools Available
		DeepSeek R1, GPT-5, Claude 4, Mythos. Full frontier ecosystem

Verizon DBIR — Vulnerability Exploitation as Initial Access Vector

DBIR Year	Data Period	Vuln Exploitation %	YoY Change	AI Context
2023	Nov 2021–Oct 2022	~5%	Baseline	No meaningful AI tooling
2024	Nov 2022–Oct 2023	14%	+180%	GPT-4 available but API-only
2025	Nov 2023–Oct 2024	20%	+34%	Claude 3.5, GPT-4o, Llama 3.1, o1
2026	Nov 2024–Oct 2025	31%	+55%	DeepSeek R1, GPT-5, Claude 4, full open-weight ecosystem

CrowdStrike — Pre-Disclosure Exploitation

Year	CVEs Exploited Before Disclosure	AI Context
2024	Not reported at this granularity	Transition period
2025	42%	Full frontier AI ecosystem available

Vulnerability Volume (Stingrai aggregation)

Year	Vulnerability Instances in Dataset
2022	68.7 million
2025	527 million (7.7× increase)

3. Breach Disclosures — The Ground Truth

If AI is genuinely accelerating attacks, mandatory breach disclosures should show it. Here's what the regulated reporting data says:

US — ITRC Annual Data Breach Report

Year	Data Compromises	YoY Change	Victim Notices
2020	~1,862	—	—
2021	~1,860	Flat	—
2022	~1,802	–3%	—
2023	3,205 (record)	+78%	353M
2024	3,158	–1%	1.35B (mega-breaches)
2025	3,322 (new record)	+5%	279M (no mega-breaches)

5-year trend: +79% increase in compromises

US — SEC Mandatory Disclosure (effective Dec 2023)

New SEC rules requiring material cyber incident disclosure within 4 business days resulted in a **60% increase in disclosures in 2024** (ITRC). However, less than 10% of notices included meaningful detail about the attack.

UK — ICO Breach Reports

Period	Cyber-Related Breaches (% of total)	Ransomware Incidents
May 2022–May 2023	25.9%	~706 (2022 full year)
May 2023–May 2024	32.5%	Increasing
2019 baseline	—	100
2020	—	440
2021	—	694
2022	—	706

UK ransomware incidents: **7× increase from 2019 to 2022** (100 → 706)

Healthcare — HIPAA/OCR (US)

Year	Large Healthcare Breaches (500+ individuals)
2018	~365

Year	Large Healthcare Breaches (500+ individuals)
2021	~714
2022	~720
2023	~745
2024	~748
2025	772 (record)

Consistent ~4% annual increase from 2022 onward, not a sudden spike.

The Opacity Problem

One critical finding undermines all breach disclosure analysis:

Year	% of Breach Notices with NO Attack Vector Information
2019	~0% (nearly all included it)
2022	40%
2023	45%
2024	65%
2025	70%

70% of breach notifications now tell you nothing about how the breach happened. This makes it structurally impossible to attribute the increase to any specific cause (AI or otherwise) from disclosure data alone.

4. The Cross-Correlation — What Lines Up and What Doesn't

What Lines Up (Correlation Supports AI Acceleration)

Vulnerability exploitation as initial access: The Verizon DBIR shows exploitation going from 5% → 14% → 20% → 31% of breaches between 2022 and 2025. This tracks almost perfectly with AI coding capability becoming available and commoditised:

- The **+180% spike (2023 data)** coincides with GPT-4's first full year of availability

- The **sustained +34% then +55% growth** coincides with open-weight models reaching exploit-capable thresholds (DeepSeek R1, Llama 3.1)
- Mandiant's TTE going negative in 2025 coincides with the full frontier AI ecosystem being available

The timing is suspicious in exactly the right way. The exploitation acceleration doesn't predate the AI tooling — it tracks it.

What Doesn't Line Up (Complicating Factors)

1. The MOVEit effect (2023): The massive +78% breach spike in 2023 was largely driven by CL0P's exploitation of MOVEit (CVE-2023-34362) — a single vulnerability, not AI-assisted, that hit thousands of organisations. This single campaign inflated 2023 numbers significantly.

2. Edge device targeting (2024-2025): Much of the vulnerability exploitation increase is concentrated in VPNs, firewalls, and edge devices. This is a **target selection shift** (attackers going after internet-facing infrastructure) not necessarily an AI capability shift. The Verizon DBIR notes edge device exploitation went from 3% to 22%.

3. Breach volume is not spiking exponentially: If AI were dramatically accelerating attacks, you'd expect breach counts to show a hockey-stick curve. Instead, ITRC shows +78% (2023), −1% (2024), +5% (2025). The growth is real but not explosive. Healthcare shows steady ~4% annual increases.

4. The disclosure opacity problem: With 70% of breach notices now omitting attack vector information, we simply cannot see whether exploitation methods have changed. The data is being deliberately obscured by reporting organisations.

5. Correlation ≠ Causation: The vulnerability exploitation increase also correlates with the massive growth in vulnerability volume (68.7M → 527M instances), the shift to remote/hybrid work exposing more edge devices, the maturation of ransomware-as-a-service ecosystems, and the MOVEit/Snowflake-style supply chain cascades. AI is not the only variable that changed.

What Would Prove It (Evidence We Don't Have Yet)

To definitively attribute the exploitation acceleration to AI, we would need:

- 1. Exploit code analysis at scale:** Systematic analysis of exploit code captured in the wild for LLM-generated signatures (the GTIG 2FA bypass is the first single case, but N=1)

2. **Attacker infrastructure analysis:** Evidence of AI model API calls or local model hosting in seized C2 infrastructure
 3. **Controlled comparison:** Exploitation speed for the same vulnerability class pre-AI vs post-AI (controlling for other variables)
 4. **Breach disclosure reform:** Mandatory attack vector reporting so we can actually see what's happening
-

5. Verdict — Stacking the Validity

The Evidence Pyramid

TIER 1 — CONFIRMED (high confidence): - Time-to-exploit has compressed from 63 days to ~7 days over 6 years (Mandiant, multi-year telemetry) - Vulnerability exploitation as initial access has gone from 5% to 31% of breaches (Verizon DBIR, 4 years of data) - AI can autonomously exploit 87% of known CVEs in lab conditions (Fang et al., replicated) - First confirmed AI-developed zero-day used in the wild (GTIG, May 2026) - Breach counts have increased ~79% over 5 years (ITRC, mandatory reporting)

TIER 2 — STRONGLY CORRELATED (moderate-high confidence): - The exploitation acceleration timeline tracks AI model availability - The cost of exploit development has collapsed (\$1/exploit, 10-15 minutes) - Nation-state actors are documented using AI for 80-90% of campaign automation - Open-weight models with no guardrails reached exploit capability in early 2025

TIER 3 — PLAUSIBLE BUT UNPROVEN (moderate confidence): - AI is the primary cause of TTE compression (vs. edge device targeting, volume growth, RaaS maturation) - Commodity criminal groups are using AI for exploitation (vs. nation-states only) - The breach disclosure increase reflects AI-accelerated attacks (vs. regulatory changes, reporting mandates)

TIER 4 — OVERSTATED (low confidence): - "All zero-days are now weaponised in hours" — only confirmed for logic flaws, not memory corruption - "AI has fundamentally changed the threat landscape for all organisations" — the evidence points to Tier 1-2 actors, not commodity criminals - "Breach numbers prove AI is accelerating attacks" — the opacity problem (70% missing attack vectors) makes this unverifiable from disclosure data

The Honest Assessment

The narrative is ahead of the ground truth, but the ground truth is catching up fast.

The AI capability is demonstrated. The exploitation is accelerating. The timelines correlate. But we cannot yet prove causation at scale from breach disclosure data because:

- 1. Breach disclosure laws don't require attack vector reporting
- 2. 70% of notifications are deliberately opaque
- 3. Multiple confounding variables changed simultaneously (edge device exposure, RaaS maturation, vulnerability volume explosion)
- 4. The single confirmed AI zero-day case (GTIG) is N=1

What we CAN say with confidence: The defender's position is deteriorating regardless of whether AI is the cause. Whether TTE went negative because of AI, because of edge device exposure, or because of RaaS maturation doesn't change the fact that it went negative. The 32-day median patch time is structurally inadequate against a 5-day (or negative) exploitation window.

The question isn't "is the hype real?" — it's "does the cause matter if the effect is the same?"

For a CISO, the answer is no. Patch faster. Detect differently. Assume breach. The attribution debate is for researchers; the operational reality is already here.

Sources

Source	Date	Data Type
Mandiant M-Trends 2025	Mar 2025	IR telemetry (500K+ hours)
Mandiant M-Trends 2026	Mar 2026	IR telemetry (500K+ hours)
Verizon DBIR 2024	May 2024	22K+ incidents
Verizon DBIR 2025	Apr 2025	22K+ incidents, 12K+ breaches
Verizon DBIR 2026	May 2026	22K+ breaches, 31K+ incidents
CrowdStrike 2026 Global Threat Report	Feb 2026	Threat intelligence
ITRC 2023 Annual Breach Report	Jan 2024	US mandatory disclosures
ITRC 2024 Annual Breach Report	Jan 2025	US mandatory disclosures

Source	Date	Data Type
ITRC 2025 Annual Breach Report	Jan 2026	US mandatory disclosures
UK ICO Data Security Incident Trends	Ongoing	UK mandatory disclosures
HIPAA Journal Healthcare Breach Stats	Jun 2026	US healthcare mandatory
Google GTIG — 2025 Zero-Day Review	Mar 2026	Threat intelligence
Google GTIG — AI Zero-Day Report	May 2026	Threat intelligence
Fang et al., UIUC	Apr 2024	Academic research
Stingrai Vulnerability Statistics 2026	Jun 2026	Aggregated statistics
Mondoo State of Vulnerabilities 2026	2026	Aggregated statistics
AI Hydra Framework (project knowledge)	Jun 2026	Internal threat model

TLP:WHITE — Analysis based on open-source intelligence, mandatory disclosure data, and project knowledge